

Deep Learning Framework: PyTorch

DR.FAKHRELDIN SAEED

1

Containt



Datasets & DataLoaders



Transforms



Automatic Differentiation with torch.autograd



Optimizing Model Parameters



Save and Load the Model



Build the Neural Network

2

PyTorch



- ▶ **PyTorch** is a Python-based machine learning library.
- ▶ It consists of two main features:
 - ▶ its ability to efficiently perform tensor operations with hardware acceleration (using GPUs)
 - ▶ and its ability to build deep neural networks.

3

Datasets & DataLoaders

- ▶ **PyTorch** provides two data primitives: (that allow you to use pre-loaded datasets as well as your own data)
 - ▶ `torch.utils.data.DataLoader`
 - ▶ `torch.utils.data.Dataset`
- ▶ **Dataset** stores the samples and their corresponding labels, and **DataLoader** wraps an iterable around the Dataset to enable easy access to the samples.
- ▶ PyTorch domain libraries provide a number of pre-loaded datasets (such as FashionMNIST) that subclass `torch.utils.data.Dataset` and implement functions specific to the particular data.
- ▶ They can be used to prototype and benchmark our model.
- ▶ We can find them here: [Image Datasets](#), [Text Datasets](#), and [Audio Datasets](#)

4

Loading a Dataset

- ▶ We load the FashionMNIST Dataset with the following parameters:
- ▶ **root** is the path where the train/test data is stored,
- ▶ **train** specifies training or test dataset,
- ▶ **download=True** downloads the data from the internet if it's not available at root.
- ▶ **transform** and **target_transform** specify the feature and label transformations

```
import torch
from torch.utils.data import Dataset
from torchvision import datasets
from torchvision.transforms import ToTensor
import matplotlib.pyplot as plt

training_data = datasets.FashionMNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor()
)

test_data = datasets.FashionMNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor()
)
```

5

Creating a Custom Dataset for your files

- ▶ A custom Dataset class must implement three functions: `__init__`, `__len__`, and `__getitem__`.
- ▶ the FashionMNIST images are stored in a directory `img_dir`, and their labels are stored separately in a CSV file `annotations_file`.
- ▶ The `__init__` function is run once when instantiating the Dataset object. We initialize the directory containing the images, the annotations file, and both transforms.
- ▶ The `__len__` function returns the number of samples in our dataset.
- ▶ The `__getitem__` function loads and returns a sample from the dataset at the given index `idx`.

```
import os
import pandas as pd
from torchvision.io import read_image

class CustomImageDataset(Dataset):
    def __init__(self, annotations_file, img_dir, transform=None, target_transform=None):
        self.img_labels = pd.read_csv(annotations_file)
        self.img_dir = img_dir
        self.transform = transform
        self.target_transform = target_transform

    def __len__(self):
        return len(self.img_labels)

    def __getitem__(self, idx):
        img_path = os.path.join(self.img_dir, self.img_labels.iloc[idx, 0])
        image = read_image(img_path)
        label = self.img_labels.iloc[idx, 1]
        if self.transform:
            image = self.transform(image)
        if self.target_transform:
            label = self.target_transform(label)
        return image, label
```

6

TRANSFORMS

- Data does not always come in its final processed form that is required for training machine learning algorithms. We use **transforms** to perform some manipulation of the data and make it suitable for training.

```
import torch
from torchvision import datasets
from torchvision.transforms import ToTensor, Lambda

ds = datasets.FashionMNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor(),
    target_transform=Lambda(lambda y: torch.zeros(10, dtype=torch.float).scatter_(0,
torch.tensor(y), value=1))
)
```

7

BUILD THE NEURAL NETWORK

```
import os
import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
```

- Get Device for Training

```
device = "cuda" if torch.cuda.is_available() else "cpu"
print(f"Using {device} device")
```

- Define the Class
- Create an instance of *NeuralNetwork*, and move it to the *device*, and print its structure.

```
model = NeuralNetwork().to(device)
print(model)
```

```
class NeuralNetwork(nn.Module):
    def __init__(self):
        super(NeuralNetwork, self).__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits
```

8

AUTOMATIC DIFFERENTIATION WITH *TORCH.AUTOGRAD*

- ▶ When training neural networks, the most frequently used algorithm is **back propagation**.
- ▶ In this algorithm, parameters (model weights) are adjusted according to the **gradient** of the loss function with respect to the given parameter.
- ▶ To compute those gradients, PyTorch has a built-in differentiation engine called ***torch.autograd***.
- ▶ It supports automatic computation of gradient for any computational graph.

```
import torch

x = torch.ones(5) # input tensor
y = torch.zeros(3) # expected output
w = torch.randn(5, 3, requires_grad=True)
b = torch.randn(3, requires_grad=True)
z = torch.matmul(x, w)+b
loss = torch.nn.functional.binary_cross_entropy_with_logits(z, y)
```

9

OPTIMIZING MODEL PARAMETERS (1/3)

- ▶ **Hyperparameters** are adjustable parameters that let you control the model optimization process. Different hyperparameter values can impact model training and convergence rates.
- ▶ hyperparameters for training:
 - ▶ **Number of Epochs** - the number times to iterate over the dataset
 - ▶ **Batch Size** - the number of data samples propagated through the network before the parameters are updated
 - ▶ **Learning Rate** - how much to update models parameters at each batch/epoch.
- ▶ Smaller values yield slow learning speed, while large values may result in unpredictable behavior during training.

```
learning_rate = 1e-3
batch_size = 64
epochs = 5
```

10

OPTIMIZING MODEL PARAMETERS (2/3)

- ▶ Once we set our hyperparameters, we can then train and optimize our model with an **optimization loop**.
- ▶ Each iteration of the optimization loop is called an **epoch**.
- ▶ Each epoch consists of two main parts:
 - ▶ The **Train Loop** - iterate over the training dataset and try to converge to optimal parameters.
 - ▶ The **Validation/Test Loop** - iterate over the test dataset to check if model performance is improving.
- ▶ **Loss function** measures the degree of dissimilarity of obtained result to the target value, and it is the loss function that we want to minimize during training.
- ▶ Common loss functions include:
 - ▶ nn.MSELoss (Mean Square Error) for regression tasks.
 - ▶ nn.NLLLoss (Negative Log Likelihood) for classification.
 - ▶ nn.CrossEntropyLoss combines nn.LogSoftmax and nn.NLLLoss.

11

OPTIMIZING MODEL PARAMETERS (3/3)

- ▶ Optimization is the process of adjusting model parameters to reduce model error in each training step.
- ▶ Optimization algorithms define how this process is performed (Stochastic Gradient Descent).
- ▶ All optimization logic is encapsulated in the optimizer object.
- ▶ There are many different optimizers available in PyTorch such as ADAM and RMSProp, that work better for different kinds of models and data.
- ▶ Inside the training loop, optimization happens in three steps:
 - ▶ Call `optimizer.zero_grad()` to reset the gradients of model parameters.
 - ▶ Backpropagate the prediction loss with a call to `loss.backward()`
 - ▶ `optimizer.step()` to adjust the parameters by the gradients collected in the backward pass.

12

Full Implementation

```
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

epochs = 10
for t in range(epochs):
    print(f"Epoch {t+1}\n-----")
    train_loop(train_dataloader, model, loss_fn, optimizer)
    test_loop(test_dataloader, model, loss_fn)
print("Done!")
```

```
def train_loop(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)
    for batch, (X, y) in enumerate(dataloader):
        # Compute prediction and loss
        pred = model(X)
        loss = loss_fn(pred, y)

        # Backpropagation
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    if batch % 100 == 0:
        loss, current = loss.item(), batch * len(X)
        print(f"loss: {loss:>7f} [{current:>5d}/{size:>5d}]")

def test_loop(dataloader, model, loss_fn):
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    test_loss, correct = 0, 0

    with torch.no_grad():
        for X, y in dataloader:
            pred = model(X)
            test_loss += loss_fn(pred, y).item()
            correct += (pred.argmax(1) == y).type(torch.float).sum().item()

    test_loss /= num_batches
    correct /= size
    print(f"Test Error: \n Accuracy: {(100*correct)/size:.1f}%, Avg loss: {test_loss:>8f} \n")
```

13

SAVE AND LOAD THE MODEL

- PyTorch models store the learned parameters in an internal state dictionary, called `state_dict`.

```
model = models.vgg16(pretrained=True)
torch.save(model.state_dict(), 'model_weights.pth')
```

- To load model weights, you need to create an instance of the same model first, and then load the parameters using `load_state_dict()` method.

```
model = models.vgg16() # we do not specify pretrained=True, i.e. do not load default weights
model.load_state_dict(torch.load('model_weights.pth'))
model.eval()
```

- Saving and Loading Models with Shapes

```
torch.save(model, 'model.pth')
```

```
model = torch.load('model.pth')
```

14