

Project: Predicting House Prices (Works as a group of 3)

Framework:

- **Problem:** Accurately predict the selling price of houses in Ames based on historical data.
- **Goal:** Develop a deep learning model using PyTorch that excels in predicting house prices, outperforming benchmark models.
- **Application:** This model can empower individuals to estimate their home's value, real estate agents to provide informed pricing recommendations, and investors to make data-driven decisions.

Data Acquisition and Exploration:

1. **Download & Import:**
 - Obtain the Ames Housing Dataset from www.kaggle.com/.
 - Use pandas to load the data into a DataFrame.
2. **Data Cleaning & Preprocessing:**
 - Handle missing values appropriately.
 - Address outliers carefully, considering domain knowledge.
 - Encode categorical features using one-hot encoding or label encoding.
 - Normalize or standardize numerical features to ensure equal importance in the model.
 - Explore feature relationships (e.g., correlations, visualizations) to gain insights.
3. **Train-Validation-Test Split:**
 - Divide the data into 70% training, 15% validation, and 15% testing sets using `sklearn.model_selection.train_test_split`.
 - This ensures robust model validation and generalizability evaluation.

Model Development and Training:

1. **Model Selection:**
 - Consider various architectures based on data characteristics and complexity:
 - **Multi-Layer Perceptron (MLP):** Good baseline for regression tasks.
 - ~~Convolutional Neural Network (CNN): Can extract spatial features if images are used.~~
 - ~~Recurrent Neural Network (RNN): Suitable for sequential data or temporal dependencies.~~
 - ~~Experiment with different architectures to find the best fit.~~

2. PyTorch Implementation:

- Start with a basic architecture like MLP, using modules like `nn.Linear`, `nn.ReLU`, and `nn.MSELoss`.
- Refer to the PyTorch tutorial (https://pytorch.org/tutorials/beginner/basics/quickstart_tutorial.html) for guidance.
- Define a forward pass to process input data and make predictions.
- Create a training loop that iterates over batches, calculates loss, uses an optimizer (e.g., Adam) to update model weights, and tracks validation performance.

3. Hyperparameter Tuning:

- Experiment with hyperparameters like learning rate, batch size, epochs, and other architecture-specific parameters.
- ~~○ Use validation set performance to guide tuning and prevent overfitting.~~
- ~~○ Consider techniques like grid search or randomized search for efficiency.~~

4. Visualization & Analysis:

- After each training session, visualize loss curves and validation metrics to track progress and identify potential issues.
- Visualize feature importance or decision boundaries to understand the model's behaviour.

Fine-tuning and Evaluation:

1. Advanced Techniques:

- If needed, explore more advanced techniques like data augmentation, regularization (dropout, L1/L2), or ensemble methods to further improve performance.

2. Benchmark Comparison:

- Compare your model's performance on the testing set with a baseline model (e.g., linear regression) and other machine learning algorithms.
- ~~○ Report metrics like MSE, R-squared, and MAE to assess accuracy and generalization.~~

~~3. Error Analysis:~~

- ~~○ Analyse errors on the testing set to identify potential biases or shortcomings in the model.~~